

Pipeline: Workflow Optimization

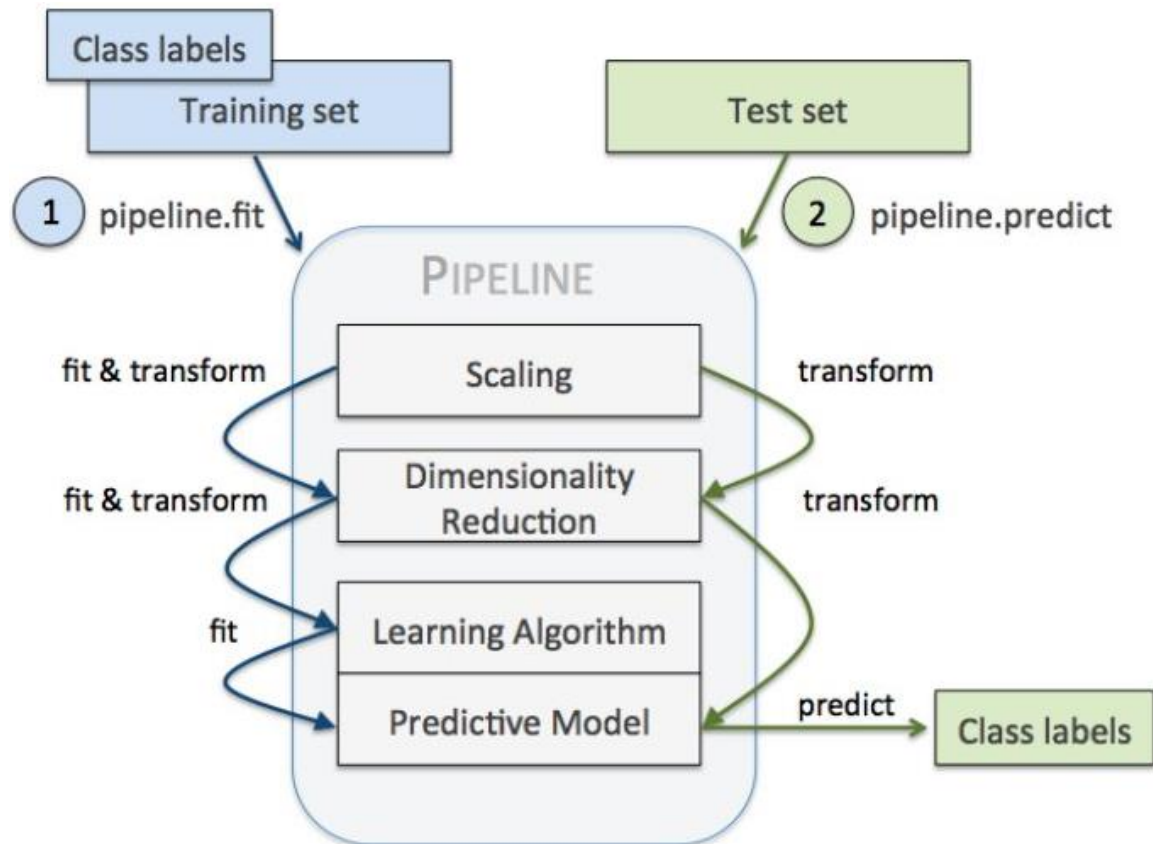


Pipeline: chaining estimators

- Pipeline can be used to chain multiple estimators into one.
- Pipeline serves two purposes:
 - Convenience and encapsulation
 - Joint parameter selection
- All estimators in a pipeline, except the last one, must be transformers.
 - The last estimator may be any type (transformer, classifier, etc.)

Pipeline: chaining estimators

- Training and prediction procedure of the pipeline



Building Pipelines

```
from sklearn.svm import SVC
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

```
# load and split the data
cancer = load_breast_cancer()
X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target, random_state=0)
```

```
from sklearn.pipeline import Pipeline
pipe = Pipeline([("scaler", MinMaxScaler()), ("svm", SVC())])
```

The **Pipeline** is built using a list of (**key**, **value**) pairs, where the **key** is a string containing the name you want to give this step and **value** is an estimator object:

```
pipe.fit(X_train, y_train).score(X_test, y_test)
```

```
0.951048951048951
```

You only have to call **fit** and **predict** once on your data to fit a whole sequence of estimators

Using Pipelines in Grid-searches

```
from sklearn.model_selection import GridSearchCV
```

```
param_grid = {'svm__C': [0.001, 0.01, 0.1, 1, 10, 100],  
              'svm__gamma': [0.001, 0.01, 0.1, 1, 10, 100]}
```

Parameters of the estimators in the pipeline should be defined using the **estimator__parameter** syntax

```
grid = GridSearchCV(pipe, param_grid=param_grid, cv=5)  
grid.fit(X_train, y_train)  
print("Best cross-validation accuracy: {:.2f}".format(  
    grid.best_score_))  
print("Test set score: {:.2f}".format(grid.score(X_test, y_test)))  
print("Best parameters: {}".format(grid.best_params_))
```

Best cross-validation accuracy: 0.98

Test set score: 0.97

Best parameters: {'svm__C': 1, 'svm__gamma': 1}

Convenient Pipeline creation with *make_pipeline*

```
from sklearn.pipeline import make_pipeline
# standard syntax
pipe_long = Pipeline([("scaler", MinMaxScaler()),
                      ("svm", SVC(C=100))])
# abbreviated syntax
pipe_short = make_pipeline(MinMaxScaler(), SVC(C=100))
```

```
print("Pipeline steps:\n{}".format(pipe_short.steps))
```


Pipeline steps:

```
[('minmaxscaler', MinMaxScaler(copy=True, feature_range=(0, 1))), ('svc',
SVC(C=100, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False))]
```

Make_pipeline does not require, and does not permit, naming the estimators. Instead, their names will be set to the **lowercase of their types** automatically.

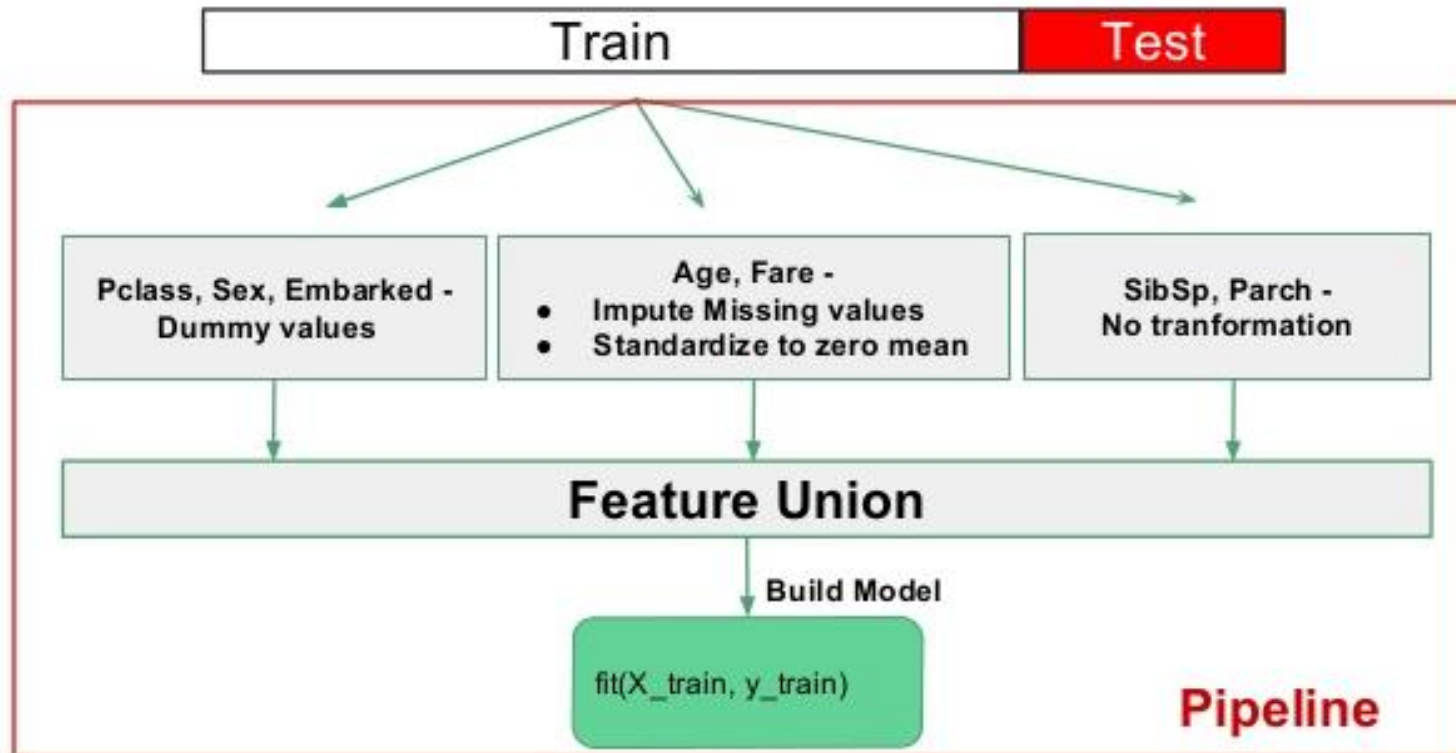
Pipeline Interface

- All estimators in a pipeline, except the last one, must be transformers. The last estimator may be any type (transformer, classifier, etc.)
- Training and prediction procedure of the pipeline



Combining Features with *FeatureUnion*

- FeatureUnion applies a list of transformer objects in parallel to the input data, then concatenates the results.
- This is useful to combine several feature extraction mechanisms into a single transformer.



Learning from Imbalanced Data



Real-life imbalanced problems

Application area	Problem description
Activity recognition	Detection of rare or less-frequent activities (multi-class problem)
Behavior analysis	Recognition of dangerous behavior (binary problem)
Cancer malignancy grading	Analyzing the cancer severity (binary and multi-class problem)
Hyperspectral data analysis	Classification of varying areas in multi-dimensional images (multi-class problem)
Industrial systems monitoring	Fault detection in industrial machinery (binary problem)
Sentiment analysis	Emotion and temper recognition in text (binary and multi-class problem)
Software defect prediction	Recognition of errors in code blocks (binary problem)
Target detection	Classification of specified targets appearing with varied frequency (multi-class problem)
Text mining	Detecting relations in literature (binary problem)
Video mining	Recognizing objects and actions in video sequences (binary and multi-class problem)

Source: "Learning from imbalanced data: open challenges and future directions", Bartosz Krawczyk



Approaches

- Sampling

- Random Undersampling
- Random Oversampling
- SMOTE
- Tomek Links
- SMOTE + Tomek Links
- Using GAN(Generative Adversarial Networks)

- Algorithms

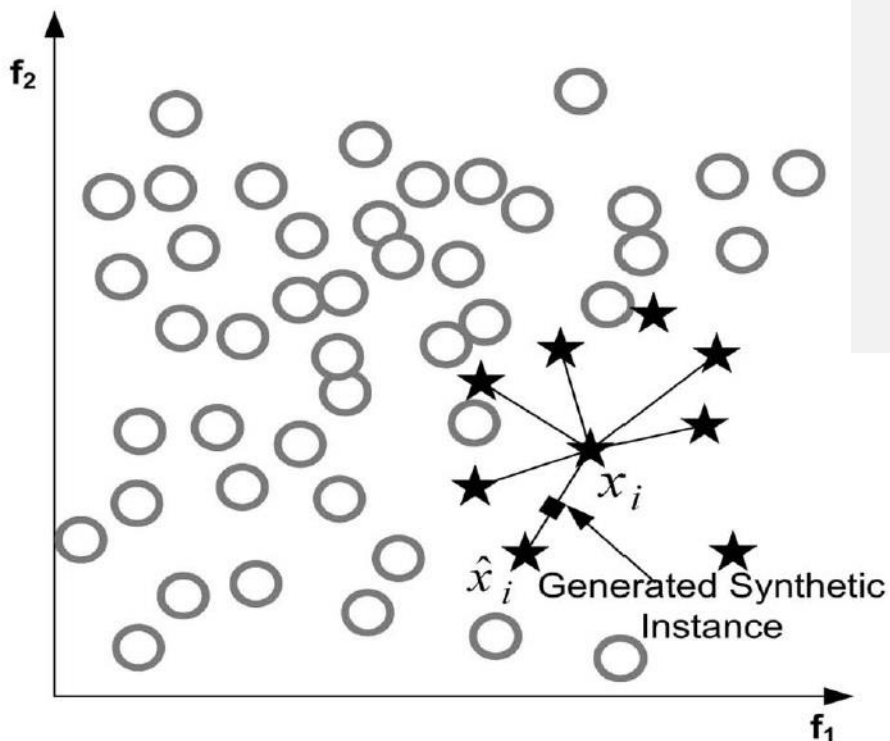
- Cost-sensitive learning methods
- Kernel-based methods



Sampling heuristics

- Consider testing under-sampling when you have a lot of data (tens- or hundreds of thousands of instances or more)
- Consider testing over-sampling when you don't have a lot of data (tens of thousands of records or less)
- Consider testing random and non-random (e.g. stratified) sampling schemes.
- Consider testing different resampled ratios (e.g. you don't have to target a 1:1 ratio in a binary classification problem, try other ratios)

SMOTE




For each point p in S (minor class):


1. Compute its k nearest neighbors in S .
2. Randomly choose $r \leq k$ of the neighbors (with replacement).
3. Choose a random point along the lines joining p and each of the r selected neighbors.
4. Add these synthetic points to the dataset with class S .

Source: "Survey of resampling techniques for improving classification performance in unbalanced datasets", Ajinkya More


Tomek links




(a)



(b)



(c)



(d)

- A pair of examples is called a ***Tomek link*** if they belong to different classes and are each other's nearest neighbors.
- Undersampling can be done by removing all tomek links from the dataset.